

소득이 50000\$ 를  
초과하는지 아닌지 분류!

(uci adult.data 활용)

# 피쳐 소개

```
In [1]: # 인구 조사 자료를 바탕으로 소득이 $ 50,000 / 년을 초과하는지 예측합니다
import pandas as pd
import numpy as np
import seaborn as sns
names=['age', 'workclass', 'fnlwgt', 'education', 'education-num', 'marital-status', 'occupation', 'relationship', 'race', 'sex', 'capital-gain', 'capital-loss', 'hours-per-week', 'native-country', 'target']
df_adult = pd.read_csv('adult.data.csv', names=names)
df_adult.head(5)
```

Out [1]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	target
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K

- age: 나이
- workclass: 노동 계급
- fnlwgt: ??
- education: 교육
- education-num: 교육 수
- marital-status: 결혼 상태
- occupation: 직업
- relationship: 가족내에서 관계
- race: 인종
- sex: 성별
- capital-gain: 자본이득
- capital-loss: 자본손실
- hours-per-week: 주당 근로시간
- native-country: 출신 국가

```
In [2]: df_adult.shape # 데이터 모양 파악
```

Out [2]: (32561, 15)

# 이상치 제거

```
In [3]: df_adult.info() # 데이터 정보 확인
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 32561 entries, 0 to 32560
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	age	32561 non-null	int64
1	workclass	32561 non-null	object
2	fnlwgt	32561 non-null	int64
3	education	32561 non-null	object
4	education-num	32561 non-null	int64
5	marital-status	32561 non-null	object
6	occupation	32561 non-null	object
7	relationship	32561 non-null	object
8	race	32561 non-null	object
9	sex	32561 non-null	object
10	capital-gain	32561 non-null	int64
11	capital-loss	32561 non-null	int64
12	hours-per-week	32561 non-null	int64
13	native-country	32561 non-null	object
14	target	32561 non-null	object

```
dtypes: int64(6), object(9)
```

```
memory usage: 3.7+ MB
```

```
In [5]: # replace()로 데이터 값 중에 ? 값 => nan 으로 바꿔주기
df_adult.replace('?', np.nan, inplace=True)
df_adult.isna().sum()
```

```
Out [5]: age                0
workclass            1836
fnlwgt                0
education             0
education-num         0
marital-status        0
occupation           1843
relationship          0
race                  0
sex                   0
capital-gain           0
capital-loss           0
hours-per-week         0
native-country        583
target                0
dtype: int64
```

```
In [6]: # dropna() 로 nan 값 없애주기
df_adult=df_adult.dropna(axis=0)
df_adult.isna().sum()
```

```
Out [6]: age                0
workclass                0
fnlwgt                    0
education                 0
education-num             0
marital-status            0
occupation                0
relationship              0
race                      0
sex                       0
capital-gain              0
capital-loss              0
hours-per-week            0
native-country            0
target                    0
dtype: int64
```

# 데이터 전처리

```
In [9]: # fnlwtg 컬럼은 필요없다고 판단하여 제거  
# marital-status, relationship 별로 상관없는 거 같은데 영향력이 커서 제거하고 해보기로함.  
df_adult.drop(['fnlwtg', 'marital-status', 'relationship'], axis=1, inplace=True)  
df_adult.shape
```

Out[9]: (30162, 12)

```
In [10]: # 레이블 인코더 사용해서 문자열 => 숫자  
from sklearn.preprocessing import LabelEncoder  
  
col_names = ['workclass', 'education', 'occupation', 'race', 'sex', 'native-country', 'target']  
le = LabelEncoder()  
for k in col_names:  
    df_adult[k] = le.fit_transform(df_adult[k])  
  
df_adult.head()
```


Out[10]:

	age	workclass	education	education-num	occupation	race	sex	capital-gain	capital-loss	hours-per-week	native-country	target
0	39	5	9	13	0	4	1	2174	0	40	38	0
1	50	4	9	13	3	4	1	0	0	13	38	0
2	38	2	11	9	5	4	1	0	0	40	38	0
3	53	2	1	7	5	2	1	0	0	40	38	0
4	28	2	9	13	9	2	0	0	0	40	4	0

```
In [11]: # 데이터 비율 맞춰주기 위해 StandardScaler 적용  
from sklearn.preprocessing import StandardScaler  
Xdf = df_adult.loc[:, 'age': 'native-country']  
ydf = df_adult.loc[:, 'target']  
  
scaler = StandardScaler()  
scaler.fit(Xdf)  
adult_scaled = scaler.transform(Xdf)
```

native-country : 0.001

<AxesSubplot :>



# 평가지표

---

```
In [16]: from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.metrics import precision_score, recall_score
from sklearn.metrics import f1_score, roc_auc_score

# 배웠던 평가지표를 구현하는 함수
def get_clf_eval(y_test , pred):
    confusion = confusion_matrix( y_test, pred)
    accuracy = accuracy_score(y_test , pred)
    precision = precision_score(y_test , pred)
    recall = recall_score(y_test , pred)
    f1 = f1_score(y_test,pred)
    roc_auc = roc_auc_score(y_test, pred)
    print('오차 행렬')
    print(confusion)
    print('정확도: {0:.4f}, 정밀도: {1:.4f}, 재현율: {2:.4f},#
    F1: {3:.4f}, AUC:{4:.4f}'.format(accuracy, precision, recall, f1, roc_auc))
```

# 결정 트리 분류

```
In [20]: #결정트리분류기
from sklearn.tree import DecisionTreeClassifier

clf=DecisionTreeClassifier(max_depth=8)
clf.fit(X_train,y_train)

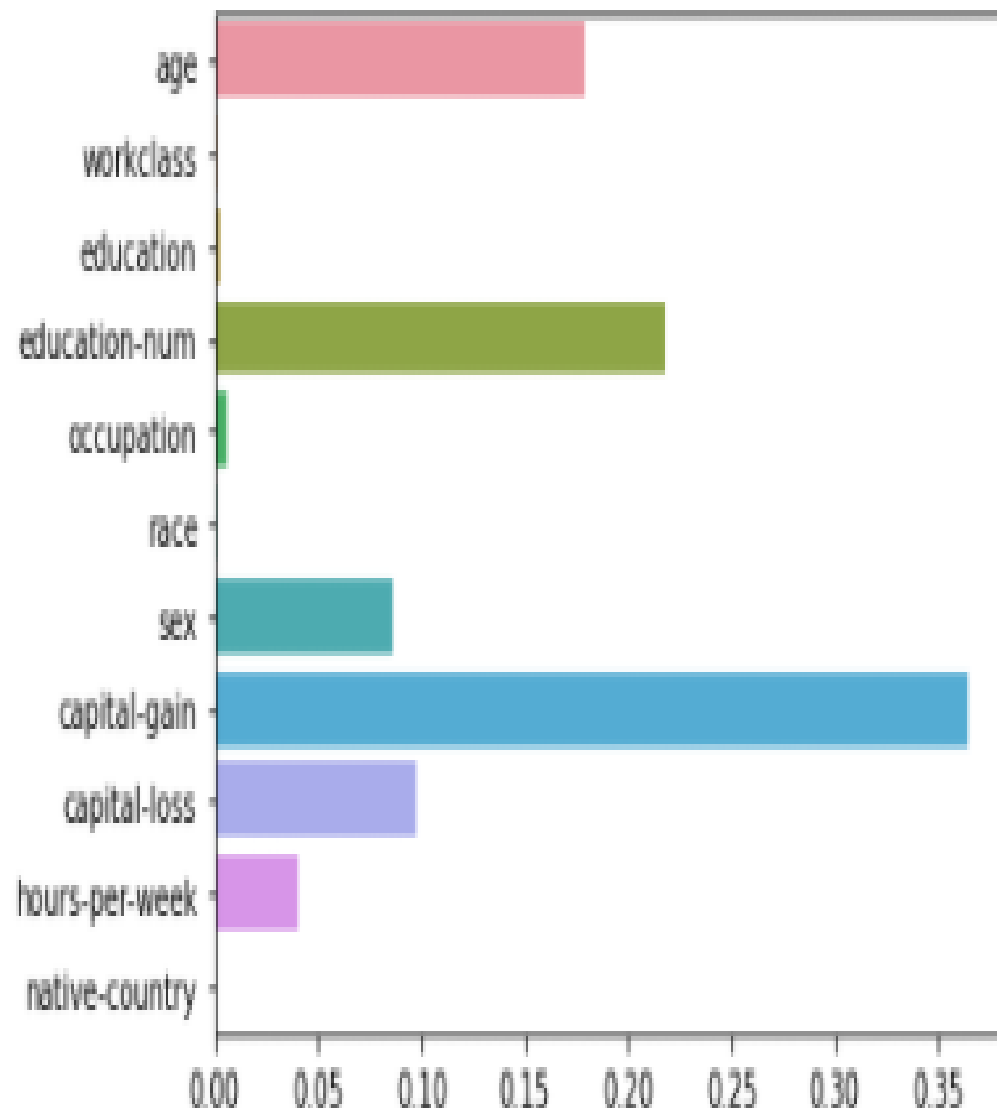
print("훈련 세트 정확도: {:.3f}".format(clf.score(X_train, y_train)))
print("테스트 세트 정확도: {:.3f}".format(clf.score(X_test, y_test)))
print()
pred=clf.predict(X_test)
get_clf_eval(y_test , pred)

훈련 세트 정확도: 0.845
테스트 세트 정확도: 0.831

오차 행렬
[[5412  227]
 [1048  854]]
정확도: 0.8309, 정밀도: 0.7900, 재현율: 0.4490, F1: 0.5726, AUC:0.7044
```

```
In [21]: import seaborn as sns
import numpy as np
%matplotlib inline
# feature importance 추출
print("Feature importances:\n{0}".format(np.round(clf.feature_importances_,
3)))
# feature별 importance 매핑
for name, value in zip(names , clf.feature_importances_):
    print('{0} : {1:.3f}'.format(name, value))
# feature importance를 column 별로 시각화 하기
sns.barplot(x=clf.feature_importances_ , y=names)
```

```
Feature importances:
[0.178 0.002 0.001 0.219 0.007 0.002 0.087 0.365 0.098 0.041 0.001]
age : 0.178
workclass : 0.002
education : 0.001
education-num : 0.219
occupation : 0.007
race : 0.002
sex : 0.087
capital-gain : 0.365
capital-loss : 0.098
hours-per-week : 0.041
native-country : 0.001
```



# 랜덤 포레스트

```
In [41]: #랜덤포레스트
from sklearn.ensemble import RandomForestClassifier

clf=RandomForestClassifier(max_depth=10, min_samples_leaf=4, min_samples_split=10, n_estimators=300)
clf.fit(X_train,y_train)


print("훈련 세트 정확도: {:.3f}".format(clf.score(X_train, y_train)))
print("테스트 세트 정확도: {:.3f}".format(clf.score(X_test, y_test)))
print()
pred=clf.predict(X_test)
get_clf_eval(y_test , pred)
```

훈련 세트 정확도: 0.849  
테스트 세트 정확도: 0.836

오차 행렬  
[[5394 245]  
 [ 995 907]]  
정확도: 0.8356, 정밀도: 0.7873, 재현율: 0.4769, F1: 0.5940, AUC:0.7167

```
In [42]: import seaborn as sns
import numpy as np
%matplotlib inline
# feature importance 추출
print("Feature importances:\n{0}".format(np.round(clf.feature_importances_,
3)))
# feature별 importance 매핑
for name, value in zip(names , clf.feature_importances_):
    print('{0} : {1:.3f}'.format(name, value))
# feature importance를 column 별로 시각화 하기
sns.barplot(x=clf.feature_importances_ , y=names)
```

Feature importances:  
[0.181 0.014 0.033 0.201 0.04 0.006 0.081 0.272 0.08 0.089 0.004]  
age : 0.181  
workclass : 0.014  
education : 0.033  
education-num : 0.201  
occupation : 0.040  
race : 0.006  
sex : 0.081  
capital-gain : 0.272  
capital-loss : 0.080  
hours-per-week : 0.089  
native-country : 0.004



# 딥러닝

```
: # 딥러닝코드
```

```
# 딥러닝을 구동하는 데 필요한 케라스 함수를 불러옵니다.
```

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import Dense
```

```
# 필요한 라이브러리를 불러옵니다.
```

```
import numpy as np
```

```
import tensorflow as tf
```

```
# 실행할 때마다 같은 결과를 출력하기 위해 설정하는 부분입니다.
```

```
np.random.seed(0)
```

```
tf.random.set_seed(0)
```

```
: from keras.callbacks import EarlyStopping
```

```
# 딥러닝 구조를 결정합니다(모델을 설정하고 실행하는 부분입니다).
```

```
model = Sequential()
```

```
model.add(Dense(30, input_dim=11, activation='relu'))
```

```
model.add(Dense(12, activation='relu'))
```

```
model.add(Dense(8, activation='relu'))
```

```
model.add(Dense(1, activation='sigmoid'))
```

```
# 딥러닝을 실행합니다.
```

```
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
early_stopping_callback=EarlyStopping(monitor='val_loss', patience=10)
```

```
model.fit(X_train, y_train, epochs=1000, batch_size=100)
```

```
Epoch 992/1000
```

```
227/227 [=====] - 0s 977us/step - loss: 0.3256 - accuracy: 0.8497
```

```
Epoch 993/1000
```

```
227/227 [=====] - 0s 749us/step - loss: 0.3265 - accuracy: 0.8503
```

```
Epoch 994/1000
```

```
227/227 [=====] - 0s 810us/step - loss: 0.3262 - accuracy: 0.8497
```

```
Epoch 995/1000
```

```
227/227 [=====] - 0s 762us/step - loss: 0.3260 - accuracy: 0.8503
```

```
Epoch 996/1000
```

```
227/227 [=====] - 0s 872us/step - loss: 0.3261 - accuracy: 0.8501
```

```
Epoch 997/1000
```

```
227/227 [=====] - 0s 806us/step - loss: 0.3260 - accuracy: 0.8497
```

```
Epoch 998/1000
```

```
227/227 [=====] - 0s 830us/step - loss: 0.3259 - accuracy: 0.8504
```

```
Epoch 999/1000
```

```
227/227 [=====] - 0s 828us/step - loss: 0.3258 - accuracy: 0.8510
```

```
Epoch 1000/1000
```

```
227/227 [=====] - 0s 819us/step - loss: 0.3260 - accuracy: 0.8504
```



	모델	정확도	정밀도	재현율	F1 score	AUC
0	딥러닝	0.8504	NaN	NaN	NaN	NaN
1	LGBM	0.8488	0.7737	0.5662	0.6539	0.7552
2	GBM	0.8472	0.7774	0.5526	0.6460	0.7496
3	XGB	0.8455	0.7744	0.5468	0.6410	0.7465
4	ada부스트	0.8373	0.7582	0.5210	0.6176	0.7325
5	랜덤포레스트	0.8356	0.7873	0.4769	0.5940	0.7167
6	결정트리	0.8309	0.7900	0.4490	0.5726	0.7044
7	svm	0.8211	0.7530	0.4327	0.5496	0.6924
8	로지스틱회귀	0.8118	0.7102	0.4290	0.5349	0.6850
9	knn	0.8088	0.6686	0.4795	0.5585	0.6997